

NovaTech Q2 2026 Product Strategy & Engineering Roadmap

1. Executive Summary & Review of Q1 Outcomes

This document establishes NovaTech's strategic engineering priorities for Q2 2026 (April 1 - June 30, 2026). During Q1 2026, the team successfully shipped the Google Drive Connector V1 and completed the Permission-Aware Retrieval Filter ahead of schedule.

Strategic Shift & Deprioritization: Following an executive review of Q1 outcomes, the Mobile Knowledge Assistant MVP was officially deprioritized and postponed from Q1 to Q4 2026. Engineering resources were reallocated entirely to RAG answer accuracy, a rigorous 3-layer evaluation framework, in-memory stream extraction, and real-time cost attribution.

2. Core Q2 Strategic Priorities and Target Deliverables

Priority 1 - 3-Layer Evaluation Benchmark Suite: Build an enterprise-grade evaluation pipeline measuring Retrieval Metrics (Recall@K, Precision@K, MRR), LLM Scorer Metrics (Groundedness, Answer Correctness, Citation Accuracy, Hallucination Rate), and Product Metrics (Query Resolution Rate, Response Latency, Feedback Satisfaction). Ships with automated markdown reporting in docs/eval-report.md. Target Completion Date: April 30, 2026.

Priority 2 - Token Usage & Cost Attribution Engine: Implement real-time token tracking per query with model-level USD cost estimation across Gemini, OpenAI, and open-source models to ensure complete financial visibility. Target Completion Date: May 15, 2026.

Priority 3 - High-Performance In-Memory Ingestion Pipeline: Build pure in-memory DOCX and PDF byte stream extraction via PyMuPDF and python-docx BytesIO, eliminating disk I/O bottlenecks and temporary file overhead during Drive and local synchronization. Target Completion Date: May 30, 2026.

Priority 4 - Phase 2 Connector Scoping: Define abstract connector interfaces for future Notion and Confluence integrations, preserving modular connector boundaries without implementing production connectors in V1. Target Completion Date: June 30, 2026.

3. Engineering Team Reallocation

Evaluation & Quality Engineering: 5 Engineers dedicated to benchmark datasets, scorer rubrics, and automated regression suites.

Core RAG Pipeline & Ingestion: 7 Engineers dedicated to in-memory extraction, ChromaDB persistence, and token tracking.

Web Application & Analytics Dashboard: 3 Engineers dedicated to Streamlit metrics dashboard.